

Accuracy — classifier (with real Context Agent enabled)

server=http://localhost:8080 classifier=v1.0-standin Context Agent: ENABLED (real LLM)
Gold: none->non_offensive; abusive/hate/violence/pornographic identical.

Comparison to the previous run WITHOUT the real Context Agent:

TEXT 5-class accuracy : 69.5 % (no-CA) -> 69.5 % (with-CA)

IMAGE 5-class accuracy : 66.3 % (no-CA) -> 66.0 % (with-CA)

=> the classifier label accuracy is unchanged by the CA (any diff is Ollama stand-in run-to-run noise). The CA changes the FINAL DECISION on escalated cases only – see the end-to-end decision page.

TEXT (/classify):

scored 200/200 (failed 0)

5-class accuracy : 69.5 % binary 79.5 % macroF1 0.490

per-class	P	R	F1	n
non_offensive	0.74	0.93	0.83	104
abusive	0.56	0.79	0.65	38
hate	1.00	0.09	0.16	23
violence	0.78	0.30	0.44	23
pornographic	0.75	0.25	0.38	12

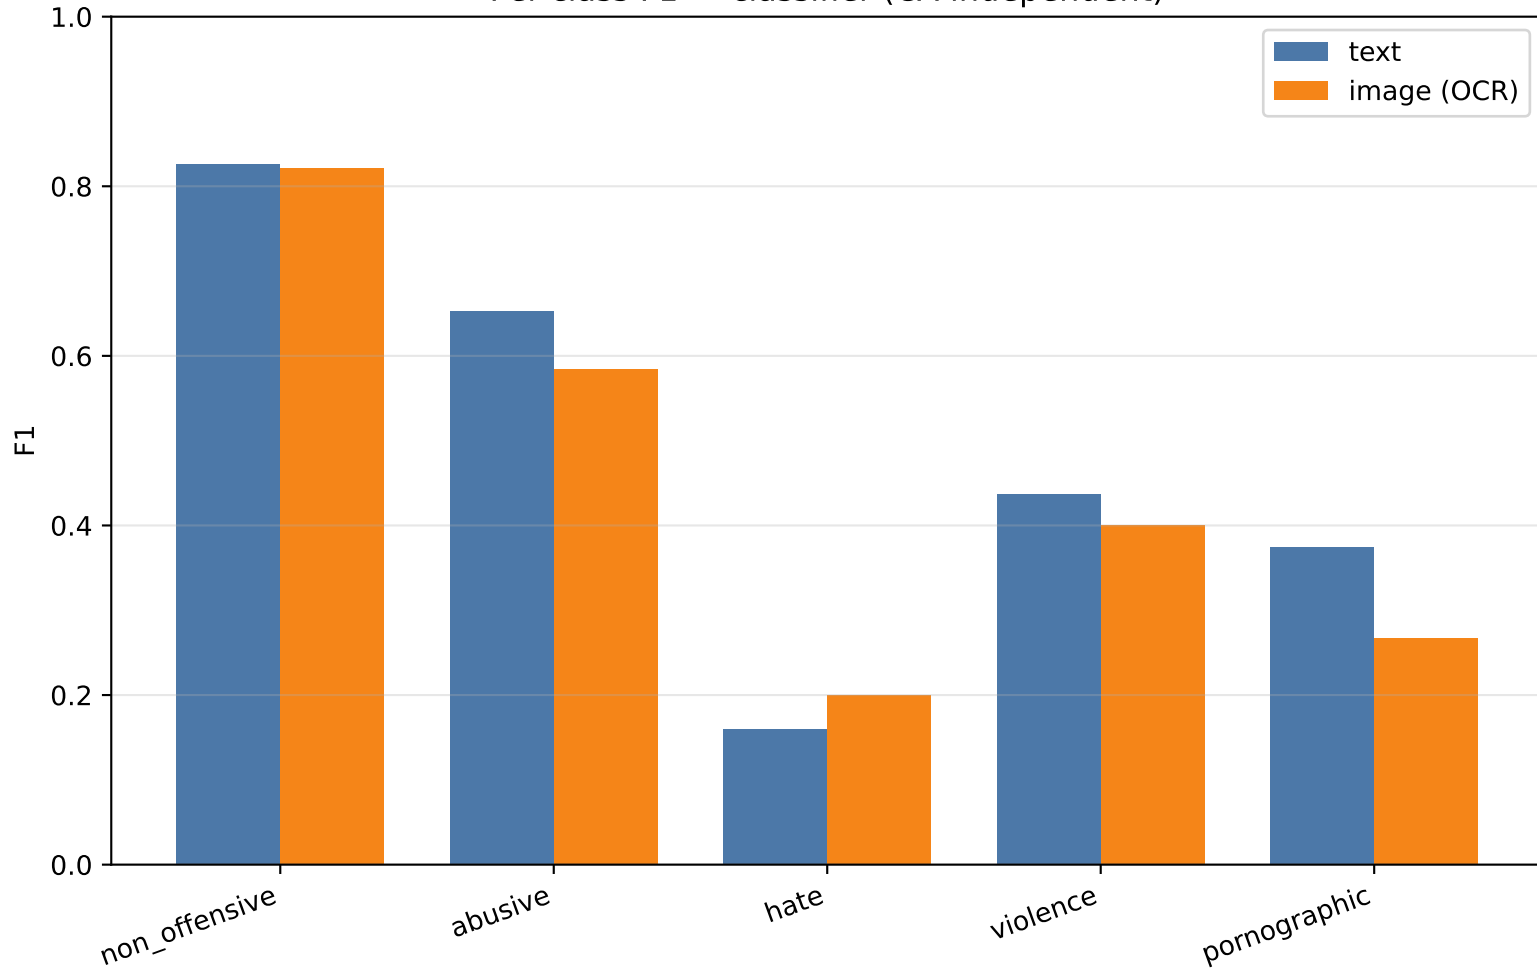
IMAGE (/classify-image):

scored 300/300 (failed 0)

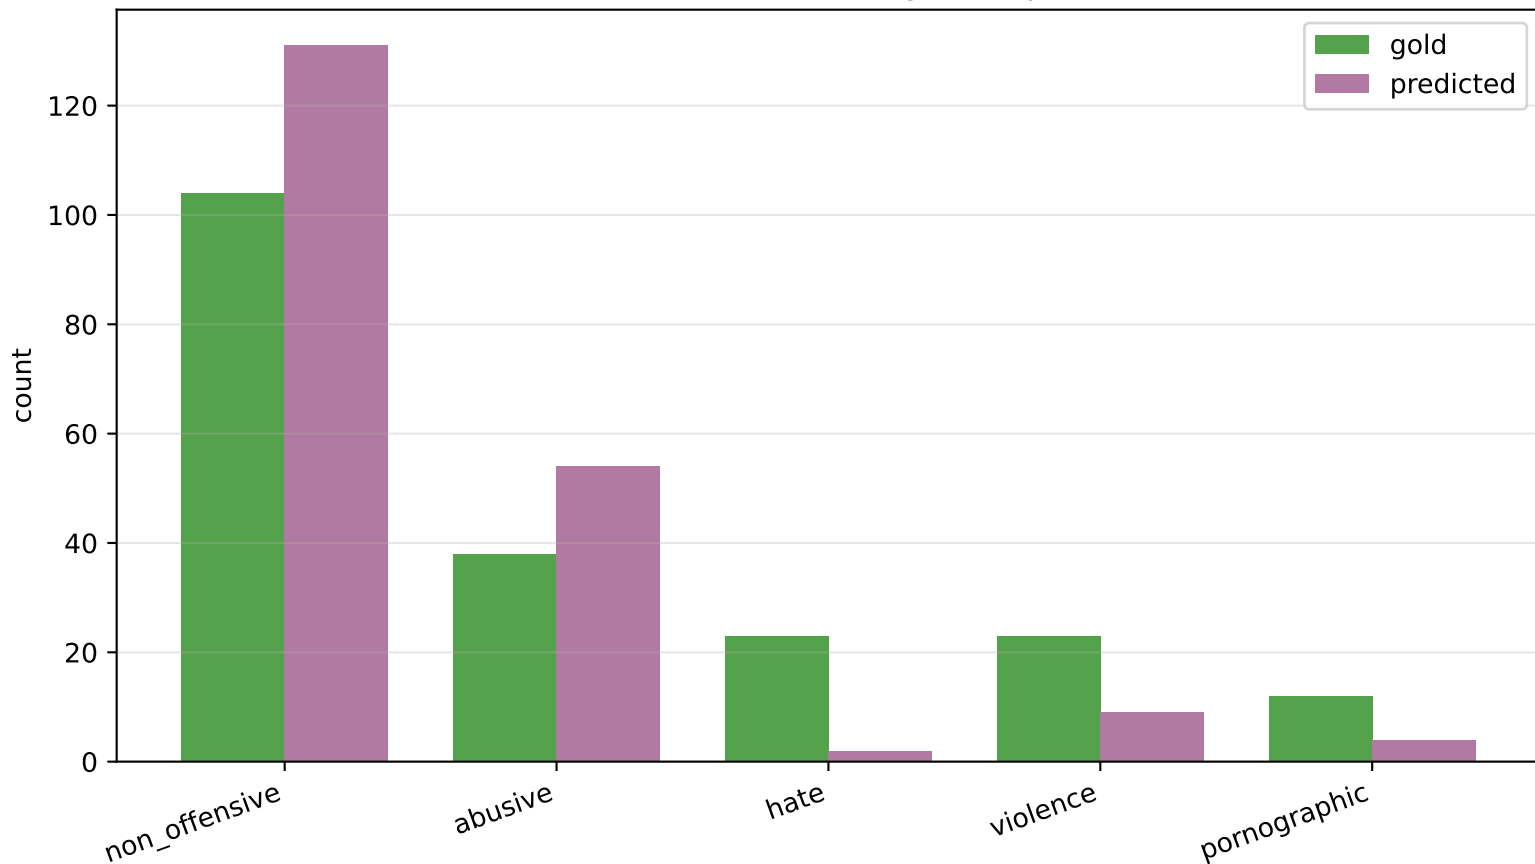
5-class accuracy : 66.0 % binary 80.3 % macroF1 0.455

per-class	P	R	F1	n
non_offensive	0.78	0.87	0.82	156
abusive	0.47	0.78	0.58	58
hate	0.80	0.11	0.20	35
violence	0.82	0.26	0.40	34
pornographic	0.31	0.24	0.27	17

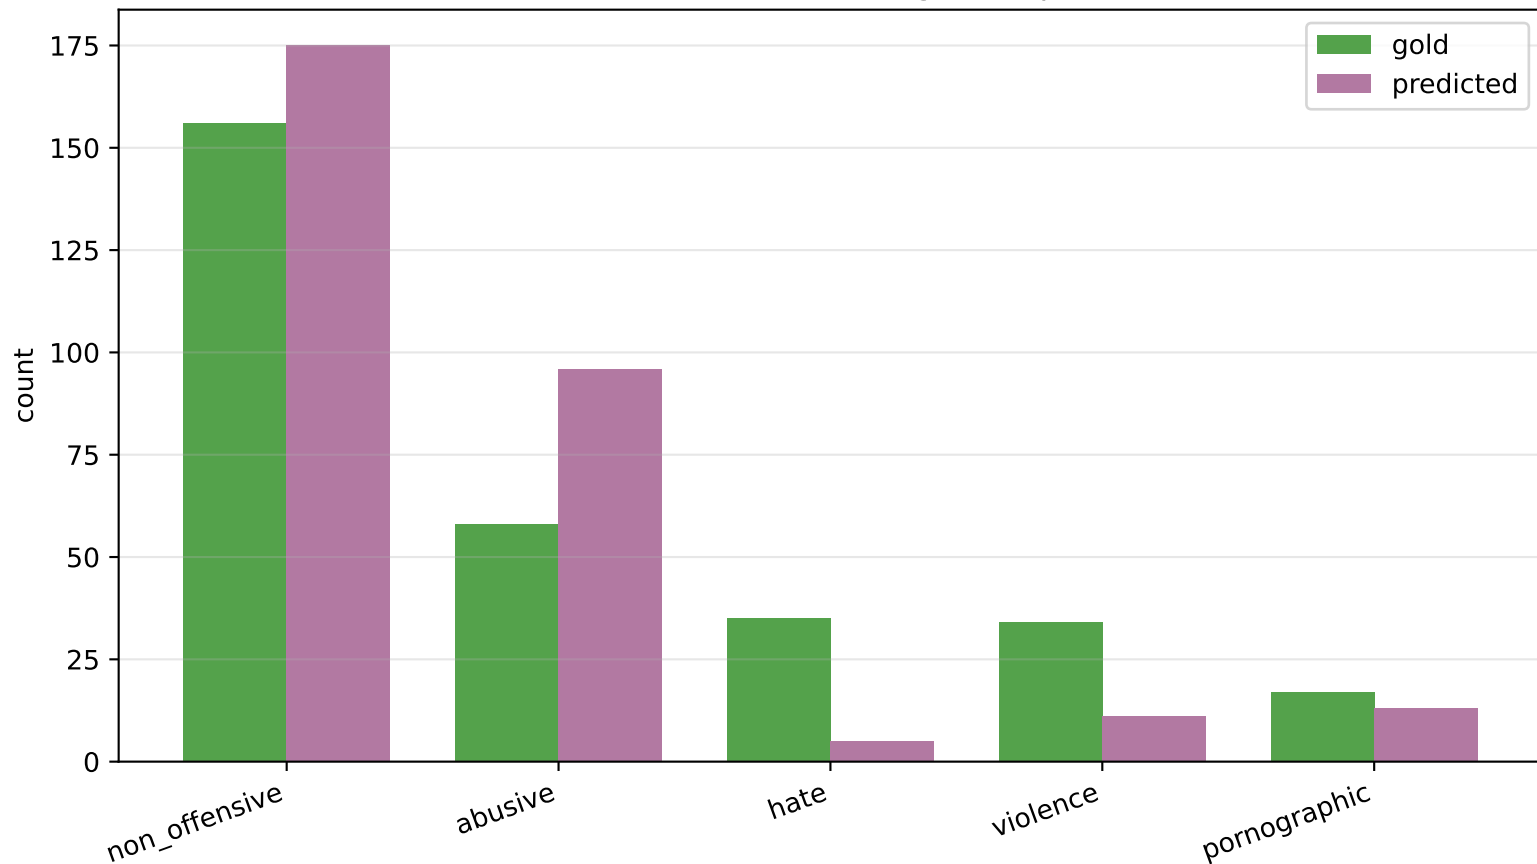
Per-class F1 — classifier (CA-independent)



Label distribution — TEXT (gold vs predicted)



Label distribution — IMAGE (gold vs predicted)



Context Agent — involvement & decision effect

The Context Agent runs ONLY on escalated cases (violence always escalates; borderline confidences would too). It does not change the classifier's label; it resolves escalated cases into ALERT / SILENT / REVIEW. 'flagged' = alert or review. recall(off)=offensive caught; FPR=non-offensive wrongly flagged.

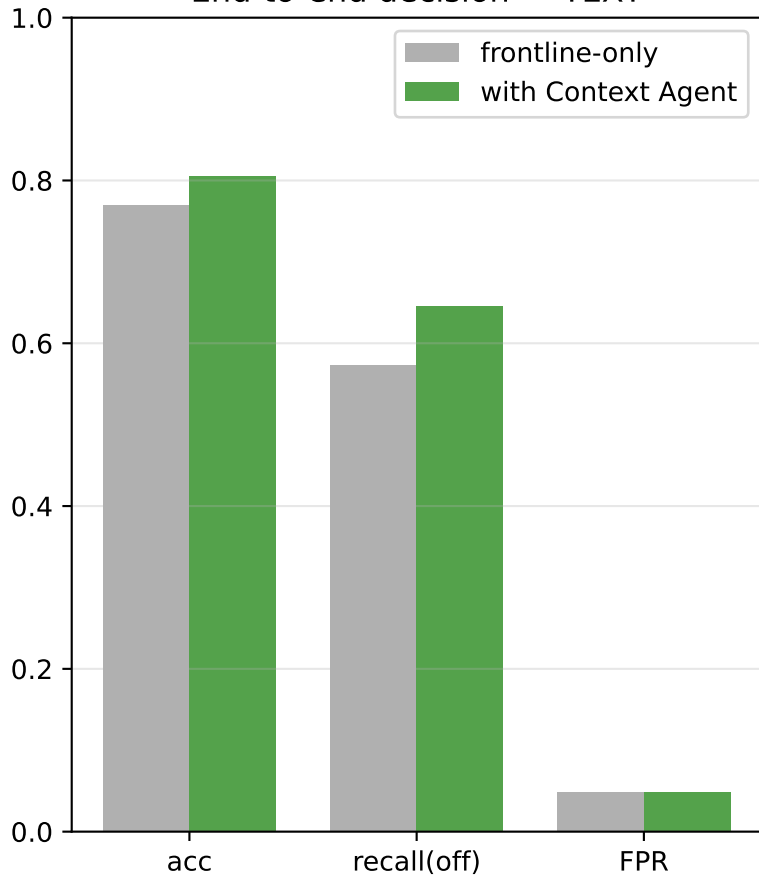
TEXT:

```
CA invoked on      : 9 samples (only escalated cases run the CA)
by LLM            : {'gemini-2.5-flash': 9}
CA verdict        : real_threat=YES 7, NO 2
escalated gold labels: {'violence': 7, 'non_offensive': 2}
end-to-end decision (n=200):
  frontline-only : acc 77.0 % recall(off) 57.3 % FPR 4.8 %
  with CA        : acc 80.5 % recall(off) 64.6 % FPR 4.8 %
delta (CA-front): acc +3.5 pp recall +7.3 pp
```

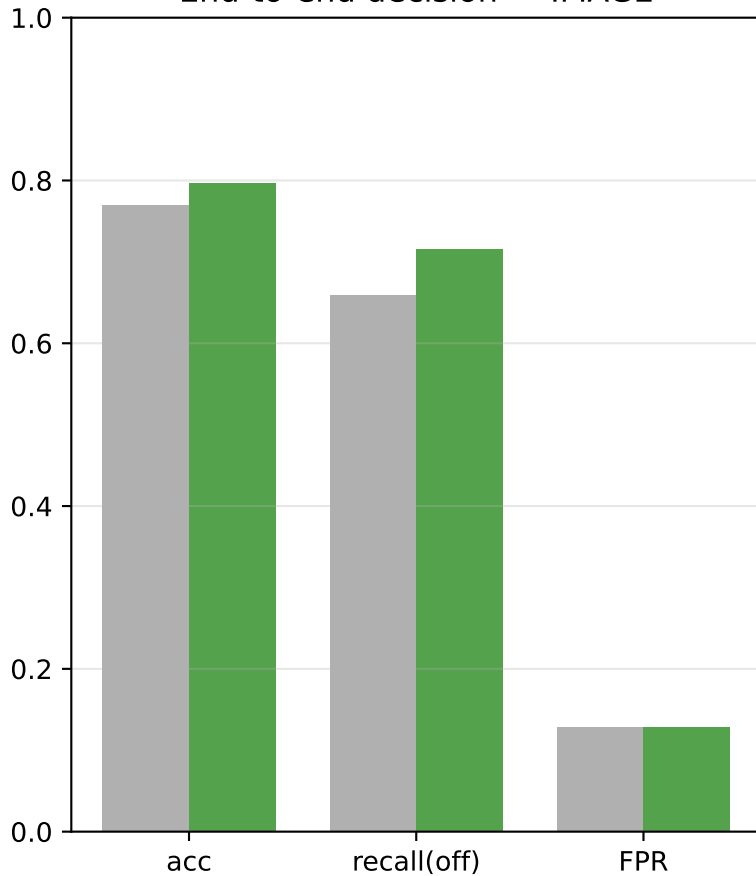
IMAGE:

```
CA invoked on      : 12 samples (only escalated cases run the CA)
by LLM            : {'gemini-2.5-flash': 12}
CA verdict        : real_threat=YES 8, NO 4
escalated gold labels: {'violence': 9, 'abusive': 1, 'non_offensive': 2}
end-to-end decision (n=300):
  frontline-only : acc 77.0 % recall(off) 66.0 % FPR 12.8 %
  with CA        : acc 79.7 % recall(off) 71.5 % FPR 12.8 %
delta (CA-front): acc +2.7 pp recall +5.6 pp
```


End-to-end decision — TEXT



End-to-end decision — IMAGE



Confusion matrix — TEXT (classifier)



Confusion matrix — IMAGE (classifier)

